

kubectl Cheat Sheet

This page contains a list of commonly used `kubectl` commands and flags.

Kubectl autocomplete

BASH

```
source <(kubectl completion bash) # setup autocomplete in bash into the current shell, bash-completion package should be installed first.
echo "source <(kubectl completion bash)" >> ~/.bashrc # add autocomplete permanently to your bash shell.
```

You can also use a shorthand alias for `kubectl` that also works with completion:

```
alias k=kubectl
complete -F __start_kubectl k
```

ZSH

```
source <(kubectl completion zsh) # setup autocomplete in zsh into the current shell
echo "[( $commands[kubectl] )] && source <(kubectl completion zsh)" >> ~/.zshrc # add autocomplete permanently to your zsh shell
```

Kubectl context and configuration

Set which Kubernetes cluster `kubectl` communicates with and modifies configuration information. See [Authenticating Across Clusters with kubeconfig](#) documentation for detailed config file information.

```

kubectl config view # Show Merged kubeconfig settings.

# use multiple kubeconfig files at the same time and view merged config
KUBECONFIG=~/.kube/config:~/.kube/kubconfig2

kubectl config view

# get the password for the e2e user
kubectl config view -o jsonpath='{.users[?(@.name == "e2e")].user.password}'

kubectl config view -o jsonpath='{.users[0].name}' # display the first user
kubectl config view -o jsonpath='{.users[*].name}' # get a list of users
kubectl config get-contexts # display list of contexts
kubectl config current-context # display the current-context
kubectl config use-context my-cluster-name # set the default context to my-cluster-name

# add a new user to your kubeconf that supports basic auth
kubectl config set-credentials kubeuser/foo.kubernetes.com --username=kubeuser --password=kubepassword

# permanently save the namespace for all subsequent kubectl commands in that context.
kubectl config set-context --current --namespace=ggckad-s2

# set a context utilizing a specific username and namespace.
kubectl config set-context gce --user=cluster-admin --namespace=foo \
&& kubectl config use-context gce

kubectl config unset users.foo # delete user foo

```

Kubectl apply

`apply` manages applications through files defining Kubernetes resources. It creates and updates resources in a cluster through running `kubectl apply`. This is the recommended way of managing Kubernetes applications on production. See [Kubectl Book](#).

Creating objects

Kubernetes manifests can be defined in YAML or JSON. The file extension `.yaml`, `.yml`, and `.json` can be used.

```

kubectl apply -f ./my-manifest.yaml           # create resource(s)
kubectl apply -f ./my1.yaml -f ./my2.yaml    # create from multiple files
kubectl apply -f ./dir                        # create resource(s) in all manifest files in dir
kubectl apply -f https://git.io/vPieo        # create resource(s) from url
kubectl create deployment nginx --image=nginx # start a single instance of nginx

# create a Job which prints "Hello World"
kubectl create job hello --image=busybox -- echo "Hello World"

# create a CronJob that prints "Hello World" every minute
kubectl create cronjob hello --image=busybox --schedule="*/1 * * * *" -- echo "Hello World"

kubectl explain pods                          # get the documentation for pod manifests

# Create multiple YAML objects from stdin
cat <<EOF | kubectl apply -f -
apiVersion: v1
kind: Pod
metadata:
  name: busybox-sleep
spec:
  containers:
  - name: busybox
    image: busybox
    args:
    - sleep
    - "1000000"
---
apiVersion: v1
kind: Pod
metadata:
  name: busybox-sleep-less
spec:
  containers:
  - name: busybox
    image: busybox
    args:
    - sleep
    - "1000"
EOF

# Create a secret with several keys
cat <<EOF | kubectl apply -f -
apiVersion: v1
kind: Secret
metadata:
  name: mysecret
type: Opaque
data:
  password: $(echo -n "s33msi4" | base64 -w0)
  username: $(echo -n "jane" | base64 -w0)
EOF

```

Viewing, finding resources

```

# Get commands with basic output
kubectl get services                # List all services in the namespace
kubectl get pods --all-namespaces  # List all pods in all namespaces

```

```

kubectll get pods -o wide           # List all pods in the current namespace, with more details
kubectll get deployment my-dep      # List a particular deployment
kubectll get pods                   # List all pods in the namespace
kubectll get pod my-pod -o yaml     # Get a pod's YAML

# Describe commands with verbose output
kubectll describe nodes my-node
kubectll describe pods my-pod

# List Services Sorted by Name
kubectll get services --sort-by=.metadata.name

# List pods Sorted by Restart Count
kubectll get pods --sort-by='.status.containerStatuses[0].restartCount'

# List PersistentVolumes sorted by capacity
kubectll get pv --sort-by=.spec.capacity.storage

# Get the version label of all pods with label app=cassandra
kubectll get pods --selector=app=cassandra -o \
  jsonpath='{.items[*].metadata.labels.version}'

# Retrieve the value of a key with dots, e.g. 'ca.crt'
kubectll get configmap myconfig \
  -o jsonpath='{.data.ca\.crt}'

# Get all worker nodes (use a selector to exclude results that have a label
# named 'node-role.kubernetes.io/master')
kubectll get node --selector='!node-role.kubernetes.io/master'

# Get all running pods in the namespace
kubectll get pods --field-selector=status.phase=Running

# Get ExternalIPs of all nodes
kubectll get nodes -o jsonpath='{.items[*].status.addresses[?(@.type=="ExternalIP")].address}'

# List Names of Pods that belong to Particular RC
# "jq" command useful for transformations that are too complex for jsonpath, it can be found at https://stedo
lan.github.io/jq/
sel=${$(kubectll get rc my-rc --output=json | jq -j '.spec.selector | to_entries | .[] | "\(.key)=\(.value),"")
}%?}
echo ${$(kubectll get pods --selector=$sel --output=jsonpath={.items..metadata.name})}

# Show labels for all pods (or any other Kubernetes object that supports labelling)
kubectll get pods --show-labels

# Check which nodes are ready
JSONPATH='{range .items[*]}{@.metadata.name}:{range @.status.conditions[*]}{@.type}={@.status}};{end}{end}' \
&& kubectll get nodes -o jsonpath="$JSONPATH" | grep "Ready=True"

# List all Secrets currently in use by a pod
kubectll get pods -o json | jq '.items[].spec.containers[].env[]?.valueFrom.secretKeyRef.name' | grep -v null
| sort | uniq

# List all containerIDs of initContainer of all pods
# Helpful when cleaning up stopped containers, while avoiding removal of initContainers.
kubectll get pods --all-namespaces -o jsonpath='{range .items[*].status.initContainerStatuses[*]}{@.containerID}
{"\n"}{end}' | cut -d/ -f3

# List Events sorted by timestamp
kubectll get events --sort-by=.metadata.creationTimestamp

# Compares the current state of the cluster against the state that the cluster would be in if the manifest wa
s applied.

```

```
kubectl diff -f ./my-manifest.yaml
```

```
# Produce a period-delimited tree of all keys returned for nodes
# Helpful when locating a key within a complex nested JSON structure
kubectl get nodes -o json | jq -c 'path(..)|[.[]|tostring]|join(".")'

# Produce a period-delimited tree of all keys returned for pods, etc
kubectl get pods -o json | jq -c 'path(..)|[.[]|tostring]|join(".")'
```

Updating resources

```
kubectl set image deployment/frontend www=image:v2 # Rolling update "www" containers of "fronte
nd" deployment, updating the image
kubectl rollout history deployment/frontend # Check the history of deployments including
the revision
kubectl rollout undo deployment/frontend # Rollback to the previous deployment
kubectl rollout undo deployment/frontend --to-revision=2 # Rollback to a specific revision
kubectl rollout status -w deployment/frontend # Watch rolling update status of "frontend"
deployment until completion
kubectl rollout restart deployment/frontend # Rolling restart of the "frontend" deployme
nt

cat pod.json | kubectl replace -f - # Replace a pod based on the JSON passed int
o std

# Force replace, delete and then re-create the resource. Will cause a service outage.
kubectl replace --force -f ./pod.json

# Create a service for a replicated nginx, which serves on port 80 and connects to the containers on port 800
0
kubectl expose rc nginx --port=80 --target-port=8000

# Update a single-container pod's image version (tag) to v4
kubectl get pod mypod -o yaml | sed 's/\(image: myimage\):.*$/\1:v4/' | kubectl replace -f -

kubectl label pods my-pod new-label=awesome # Add a Label
kubectl annotate pods my-pod icon-url=http://goo.gl/XXBTWq # Add an annotation
kubectl autoscale deployment foo --min=2 --max=10 # Auto scale a deployment "foo"
```

Patching resources

```

# Partially update a node
kubectl patch node k8s-node-1 -p '{"spec":{"unschedulable":true}}'

# Update a container's image; spec.containers[*].name is required because it's a merge key
kubectl patch pod valid-pod -p '{"spec":{"containers":[{"name":"kubernetes-serve-hostname","image":"new image"}]}}'

# Update a container's image using a json patch with positional arrays
kubectl patch pod valid-pod --type='json' -p='[{"op": "replace", "path": "/spec/containers/0/image", "value": "new image"}]'

# Disable a deployment livenessProbe using a json patch with positional arrays
kubectl patch deployment valid-deployment --type json -p='[{"op": "remove", "path": "/spec/template/spec/containers/0/livenessProbe"}]'

# Add a new element to a positional array
kubectl patch sa default --type='json' -p='[{"op": "add", "path": "/secrets/1", "value": {"name": "whatever"}}]'

```

Editing resources

Edit any API resource in your preferred editor.

```

kubectl edit svc/docker-registry          # Edit the service named docker-registry
KUBE_EDITOR="nano" kubectl edit svc/docker-registry  # Use an alternative editor

```

Scaling resources

```

kubectl scale --replicas=3 rs/foo          # Scale a replicaset named 'foo' to 3
kubectl scale --replicas=3 -f foo.yaml     # Scale a resource specified in "foo.yaml"
to 3
kubectl scale --current-replicas=2 --replicas=3 deployment/mysql # If the deployment named mysql's current size is 2, scale mysql to 3
kubectl scale --replicas=5 rc/foo rc/bar rc/baz # Scale multiple replication controllers

```

Deleting resources

```

kubectl delete -f ./pod.json              # Delete a pod using the type and name specified in pod.json
kubectl delete pod,service baz foo        # Delete pods and services with same names "baz" and "foo"
kubectl delete pods,services -l name=myLabel # Delete pods and services with label name=myLabel
kubectl -n my-ns delete pod,svc --all      # Delete all pods and services in namespace my-ns,
# Delete all pods matching the awk pattern1 or pattern2
kubectl get pods -n mynamespace --no-headers=true | awk '/pattern1|pattern2/{print $1}' | xargs kubectl delete -n mynamespace pod

```

Interacting with running Pods

kubectl logs my-pod	# dump pod logs (stdout)
kubectl logs -l name=myLabel	# dump pod logs, with label name=myLabel (stdout)
kubectl logs my-pod --previous	# dump pod logs (stdout) for a previous instantiation of a container
kubectl logs my-pod -c my-container	# dump pod container logs (stdout, multi-container case)
kubectl logs -l name=myLabel -c my-container	# dump pod logs, with label name=myLabel (stdout)
kubectl logs my-pod -c my-container --previous	# dump pod container logs (stdout, multi-container case) for a previous instantiation of a container
kubectl logs -f my-pod	# stream pod logs (stdout)
kubectl logs -f my-pod -c my-container	# stream pod container logs (stdout, multi-container case)
)	
kubectl logs -f -l name=myLabel --all-containers	# stream all pods logs with label name=myLabel (stdout)
kubectl run -i --tty busybox --image=busybox -- sh	# Run pod as interactive shell
kubectl run nginx --image=nginx -n mynamespace	# Run pod nginx in a specific namespace
kubectl run nginx --image=nginx .yaml	# Run pod nginx and write its spec into a file called pod.yaml
--dry-run=client -o yaml > pod.yaml	
kubectl attach my-pod -i	# Attach to Running Container
kubectl port-forward my-pod 5000:6000	# Listen on port 5000 on the local machine and forward to port 6000 on my-pod
kubectl exec my-pod -- ls /	# Run command in existing pod (1 container case)
kubectl exec --stdin --tty my-pod -- /bin/sh	# Interactive shell access to a running pod (1 container case)
kubectl exec my-pod -c my-container -- ls /	# Run command in existing pod (multi-container case)
kubectl top pod POD_NAME --containers	# Show metrics for a given pod and its containers

Interacting with Nodes and cluster

kubectl cordon my-node	# Mark my-node as unschedulable
kubectl drain my-node	# Drain my-node in preparation for maintenance
tenance	
kubectl uncordon my-node	# Mark my-node as schedulable
kubectl top node my-node	# Show metrics for a given node
kubectl cluster-info	# Display addresses of the master and services
kubectl cluster-info dump	# Dump current cluster state to stdout
kubectl cluster-info dump --output-directory=/path/to/cluster-state	# Dump current cluster state to /path/to/cluster-state
# If a taint with that key and effect already exists, its value is replaced as specified.	
kubectl taint nodes foo dedicated=special-user:NoSchedule	

Resource types

List all supported resource types along with their shortnames, [API group](#), whether they are [namespaced](#), and [Kind](#):

```
kubectl api-resources
```

Other operations for exploring API resources:

```
kubectl api-resources --namespaced=true      # All namespaced resources
kubectl api-resources --namespaced=false     # All non-namespaced resources
kubectl api-resources -o name                # All resources with simple output (just the resource name)
kubectl api-resources -o wide                # All resources with expanded (aka "wide") output
kubectl api-resources --verbs=list,get       # All resources that support the "list" and "get" request verbs
kubectl api-resources --api-group=extensions # All resources in the "extensions" API group
```

Formatting output

To output details to your terminal window in a specific format, add the `-o` (or `--output`) flag to a supported `kubectl` command.

Output format	Description
<code>-o=custom-columns=<spec></code>	Print a table using a comma separated list of custom columns
<code>-o=custom-columns-file=<filename></code>	Print a table using the custom columns template in the <code><filename></code> file
<code>-o=json</code>	Output a JSON formatted API object
<code>-o=jsonpath=<template></code>	Print the fields defined in a jsonpath expression
<code>-o=jsonpath-file=<filename></code>	Print the fields defined by the jsonpath expression in the <code><filename></code> file
<code>-o=name</code>	Print only the resource name and nothing else
<code>-o=wide</code>	Output in the plain-text format with any additional information, and for pods, the node name is included
<code>-o=yaml</code>	Output a YAML formatted API object

Examples using `-o=custom-columns`:

```
# All images running in a cluster
kubectl get pods -A -o=custom-columns='DATA:spec.containers[*].image'

# All images excluding "k8s.gcr.io/coredns:1.6.2"
kubectl get pods -A -o=custom-columns='DATA:spec.containers[?(@.image!="k8s.gcr.io/coredns:1.6.2")].image'

# All fields under metadata regardless of name
kubectl get pods -A -o=custom-columns='DATA:metadata.*'
```

More examples in the `kubectl` [reference documentation](#).

Kubectl output verbosity and debugging

Kubectl verbosity is controlled with the `-v` or `--v` flags followed by an integer representing the log level. General Kubernetes logging conventions and the associated log levels are described [here](#).

Verbosity	Description
<code>--v=0</code>	Generally useful for this to <i>always</i> be visible to a cluster operator.

<code>--v=1</code>	A reasonable default log level if you don't want verbosity.
<code>--v=2</code>	Useful steady state information about the service and important log messages that may correlate to significant changes in the system. This is the recommended default log level for most systems.
<code>--v=3</code>	Extended information about changes.
<code>--v=4</code>	Debug level verbosity.
<code>--v=6</code>	Display requested resources.
<code>--v=7</code>	Display HTTP request headers.
<code>--v=8</code>	Display HTTP request contents.
<code>--v=9</code>	Display HTTP request contents without truncation of contents.

What's next

- Read the [kubect! overview](#) and learn about [JsonPath](#).
- See [kubect!](#) options.
- Also read [kubect! Usage Conventions](#) to understand how to use kubect! in reusable scripts.
- See more community [kubect! cheatsheets](#).

Feedback

Was this page helpful?